

Docker

The Complete Guide

with

Hands-on Tutorial

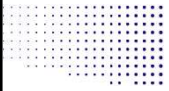
Presenter: Bach Minh Nam

Oct 2020



Agenda

1. Docker
2. Docker compose
3. Docker Swarm



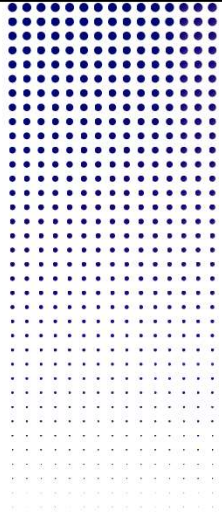
Agenda

1. Docker

1. *What is Docker? Why do we need Docker?*
2. *Basic concepts: Docker Engine, container, image, registry, basic commands*
3. *Core concepts of Container*
4. *How to build my own image? Dockerfile AZ*

2. Docker compose

3. Docker Swarm



Docker

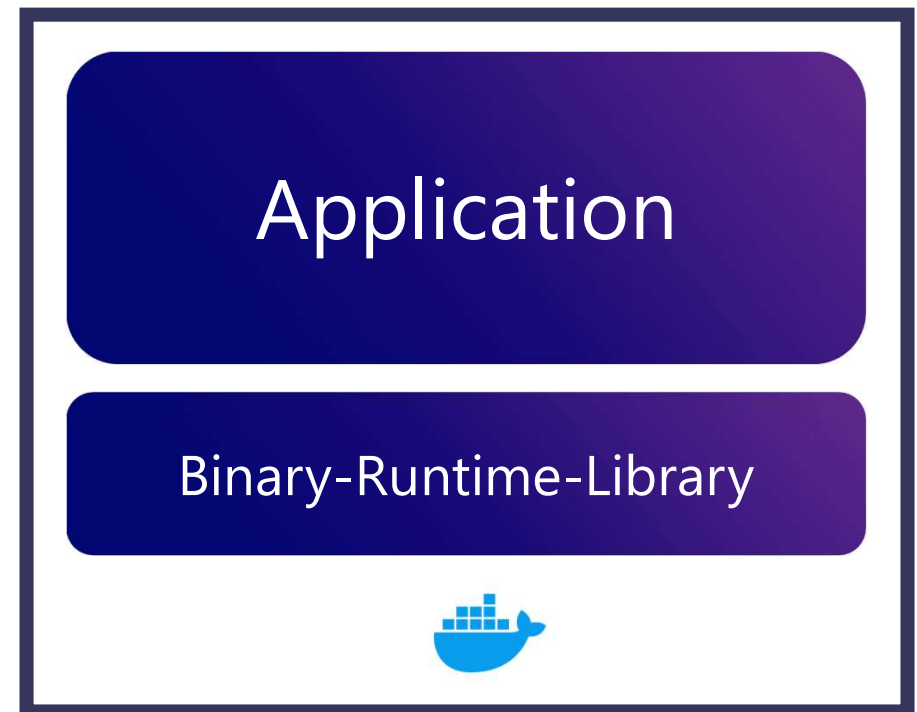
1

1. What is Docker? Why do we need Docker?

1. What is Docker?

When developing application, it usually needs *belonging dependencies* such as binary files, libraries, runtime environment...

Docker is a technology allows us to wrap the *application* and its *dependencies* into **one package**, which are **portable** (run anywhere) and **executable** (run anytime).



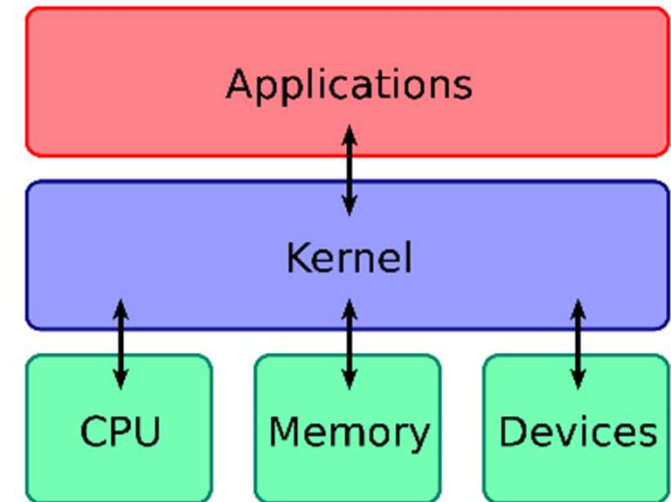
1. What is Docker: *container, image*

"Docker is a technology allows us to wrap the *application* and its *dependencies* into **one package**, which are **portable** (run anywhere) and **executable** (run anytime)"

- Those packages are called **images**.
- When executing an image, we get a **container**.
- Containers functions like a virtual machine with fully provided features such as file system, network interface, process tree...
- However, containers are not virtual machines at all. They don't have their own OS kernel, but **share the same kernel** with the physical machine.
- In the end, they are just ***processes*** running on the OS and managed by Docker.

1. What is Docker: OS kernel

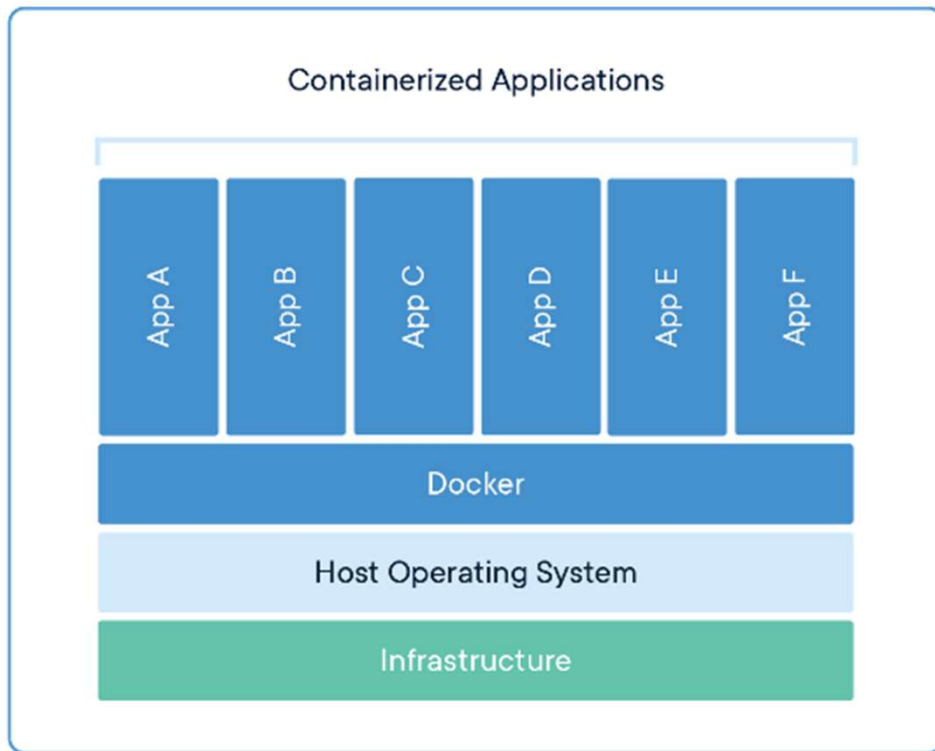
- Kernel is the core program, the heart of the operating system.
- It controls over everything in the system.
- It facilitates interactions between hardware and software components



1. What is Docker: *container vs VM*

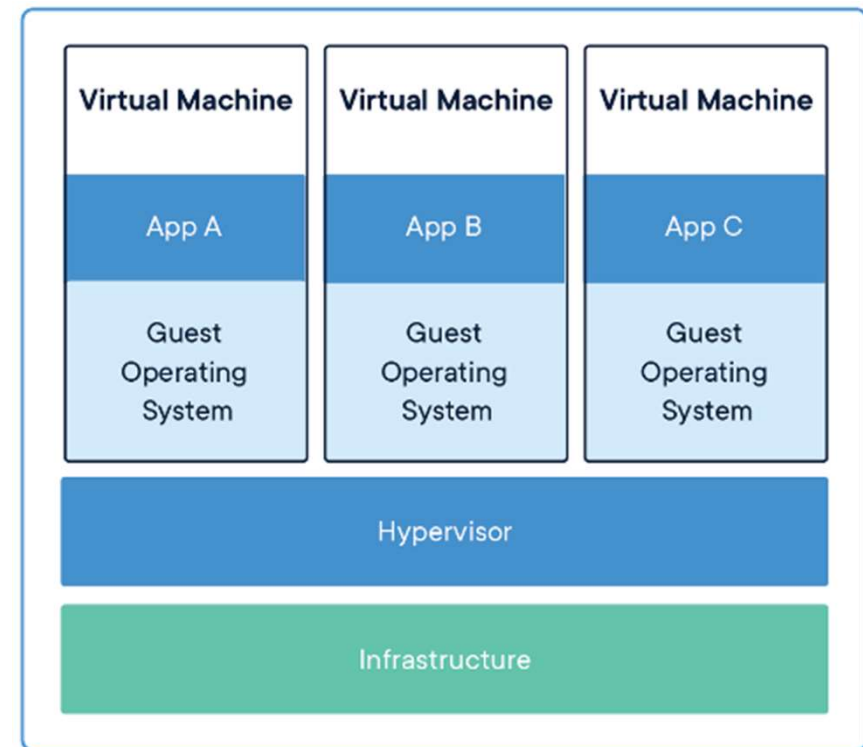
Container

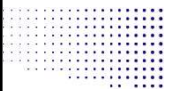
Abstraction at *software* layer - share same OS kernel
Lightweight (megabytes)



Virtual Machine

Abstraction at *hardware* layer – run its own OS
Heavy (gigabytes)





1. What is Docker: underlying technology

Docker is written in the **Go** programming language and takes advantage of several features of the Linux kernel to deliver its functionality:

- Namespaces: pid, net, ipc, mnt, uts
- Control groups: hardware resources
- Union file systems: layers creation
- Container format: format of Docker image

1. Why do we need Docker?

Docker helps to package (containerize) and ship and run application more easily.

Manually



With Docker

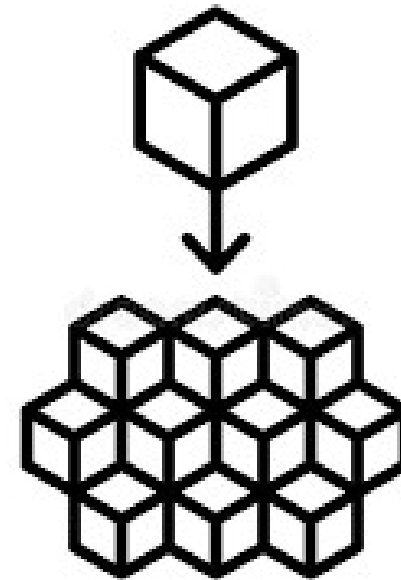


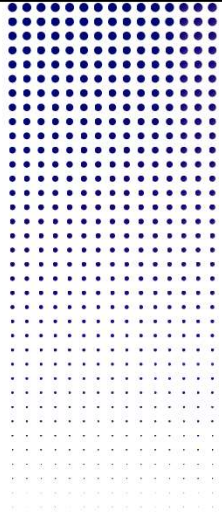
1. Why do we need Docker?

Docker helps to package (containerize) and ship and run application more easily.

Microservices

- *Immutable*: same behavior
- *Lightweight*: fast creation
- *Stateless*: disposable and ephemeral





Docker

1

2. Basic concepts: Docker Engine, container, image, registry

2. Basic concepts: *Docker Engine*

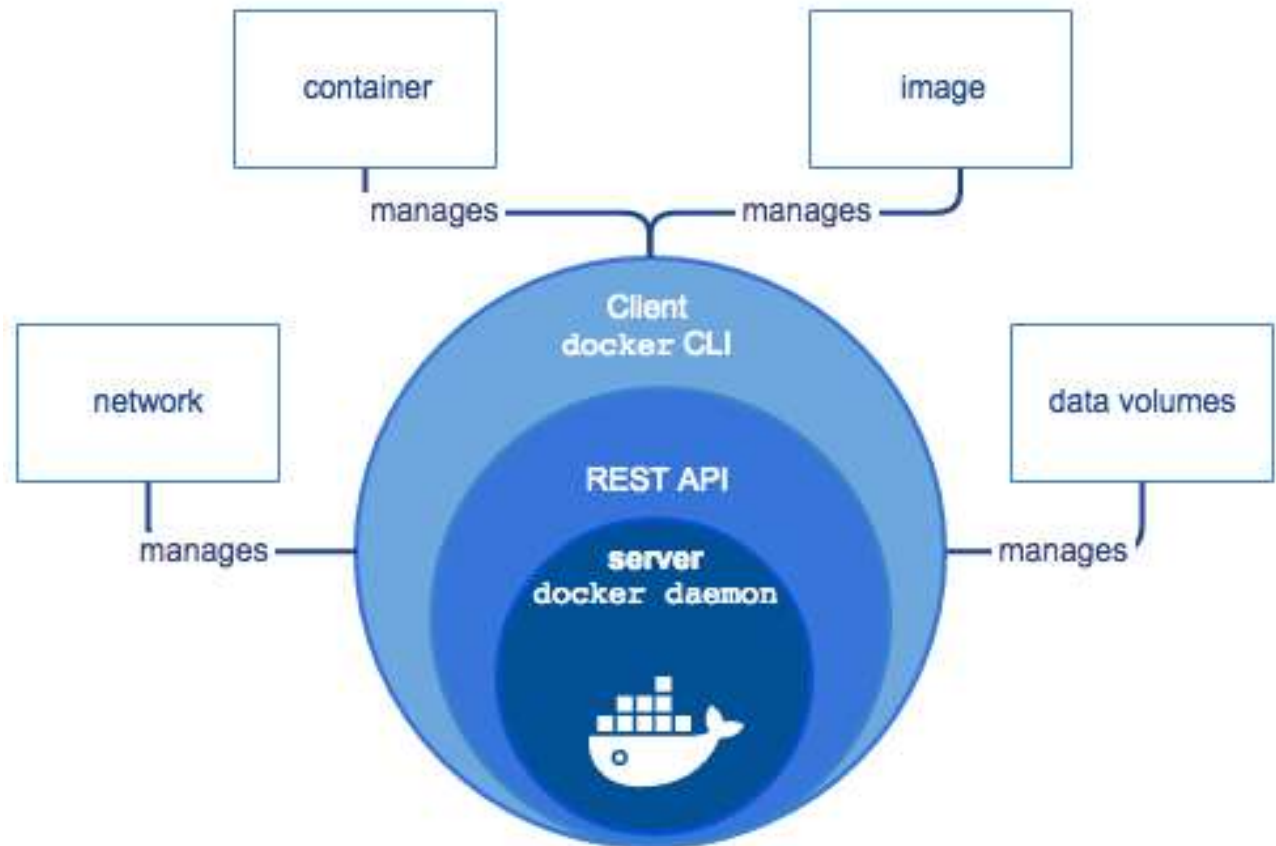
Docker Engine

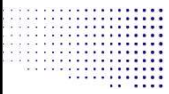
- Docker Daemon & exposed APIs
- Docker client (cli)

Share OS kernel

⇒ OS dependent

⇒ On Windows must install a Linux virtual machine to run Linux containers

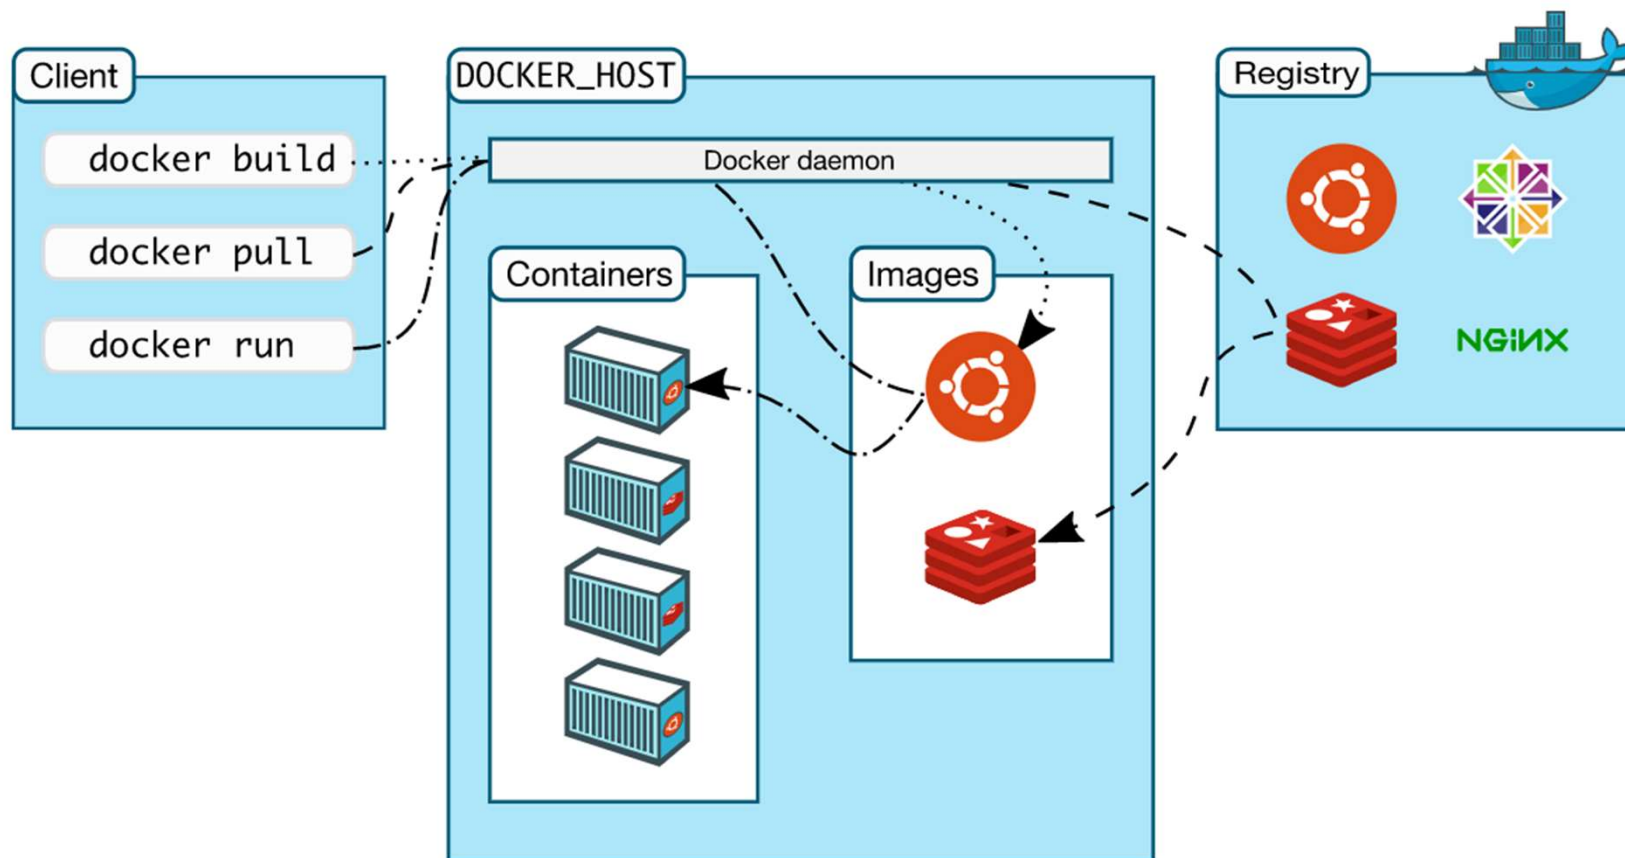


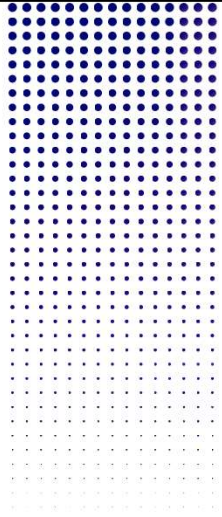


2. Basic concepts : *registry*

- Docker packages (containerizes) application and its dependencies into **image**.
- Images are stored in Docker repository call **container registry**.
- Some Docker registry providers:
 - Docker Hub
 - Amazon Elastic Container Registry (ECR)
 - Google Container Registry (GCR)
 - Azure Container Registry (ACR)

2. Basic concepts: *overall*





Docker

1

2. Basic commands: *demo*

2. Basic commands

Syntax

docker <component> <command>

Components:

- image
- container
- network
- volume
- ...

Commands:

- ls: list
- run
- exec
- stop
- pull
- prune
- ...



2. Basic commands: *image*

```
docker image pull <image>
docker image pull <image>:<tag>
docker image push <image>:<tag>
docker image ls | docker images
docker image prune
```

Short-hand:

```
docker pull
docker push
```



2. Basic commands: *container*

```
docker container run <image>
```

```
docker container ls | docker container ls -a
```

```
docker ps | docker ps -a
```

```
docker container stop <container_id>
```

```
docker container prune
```

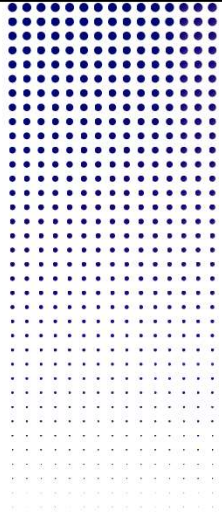
```
docker container exec <container_id> <command>
```

Short-hand:

```
docker run
```

```
docker stop
```

```
docker exec
```

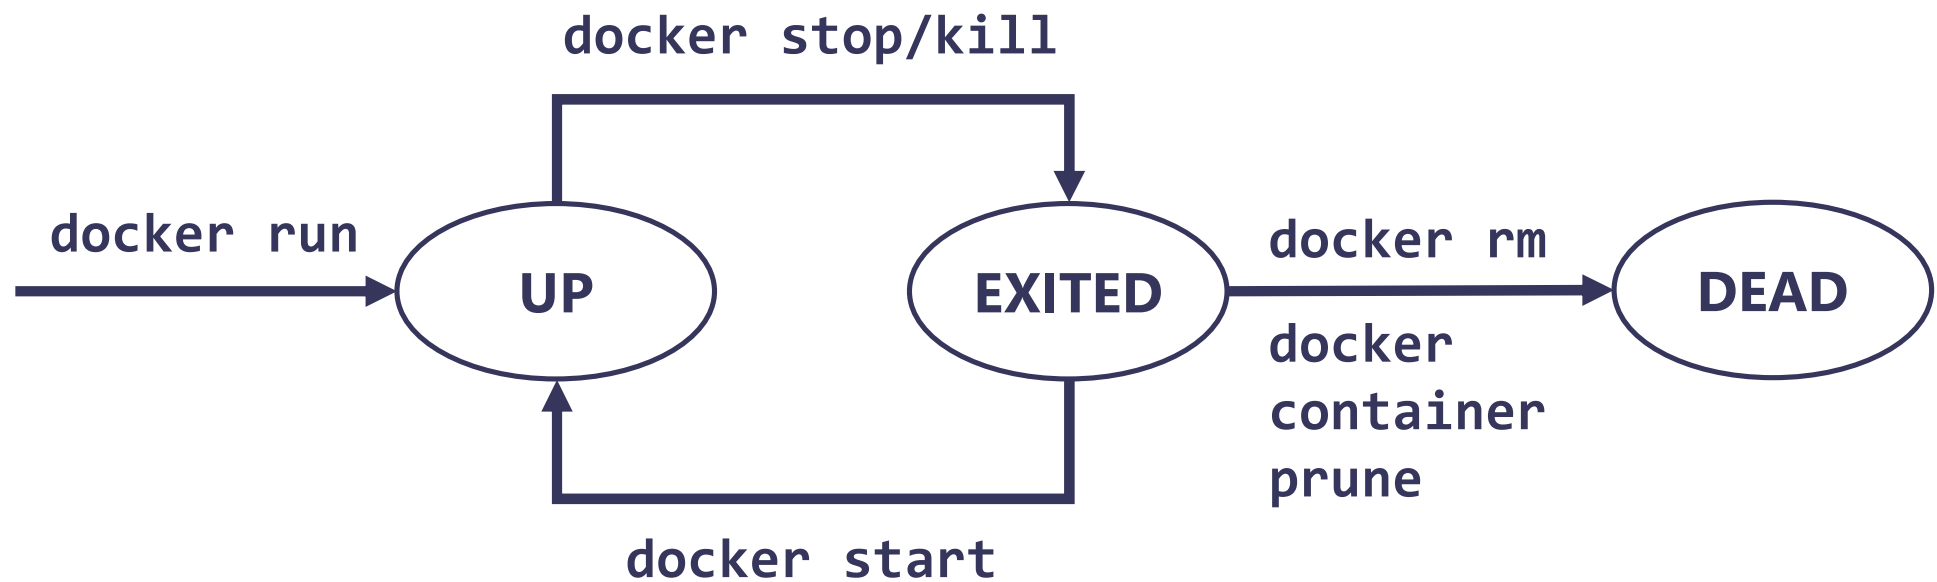


Docker

1

3. Core concepts of Container

3. Container Life Cycle





3. Container Life Cycle in Depth: PID 1

- Life cycle of a container depends on life cycle of the process running inside that has **PID = 1** (main process). The container will remain alive *as long as* the PID 1 process is alive.
- When using docker stop, Docker will send **SIGTERM** (terminating signal) to stop process PID 1 inside, then the container will stop with **EXITED (0)**.
- Within *10 seconds*, if container does not exit, Docker will send **SIGKILL** (kill signal) and the container will stop immediately with **EXITED (137)**.

3. Container Life Cycle in Depth: PID 1

`docker run alpine`

- Why does the container exit immediately?

`docker run -it alpine`

- `-i` or `--interactive`: keep STDIN open even if not attached
- `-t` or `--tty`: allocate a pseudo-TTY (TeleTYpewriter)
- To exit the shell within a container, press **Ctrl + D** because Ctrl + C does not work.
- Exit the shell also makes the container exit. How to keep the container alive?
- Attach – detach mode: press **Ctrl + P + Q** to detach container from local terminal; `docker attach <container_id>` to re-attach container stdin and stdout into local terminal.

`docker run -d <image>`: to run a container in detach mode



3. Execute commands inside container

Syntax

- `docker exec <container_id> <command>`

Example

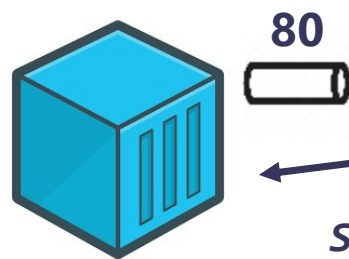
- `docker exec <container_id> echo Hello World`
- `docker exec <container_id> echo $PATH`
- `docker exec <container_id> sh -c "echo $PATH"`
- `docker exec <container_id> cat /etc/os-release`

Good To Remember

To SSH into any container, use `exec` shell or `bash` with `-it`

- `docker exec -it <container_id> sh`
- `docker exec -it <container_id> bash`

3. Port Mapping



nginx

80

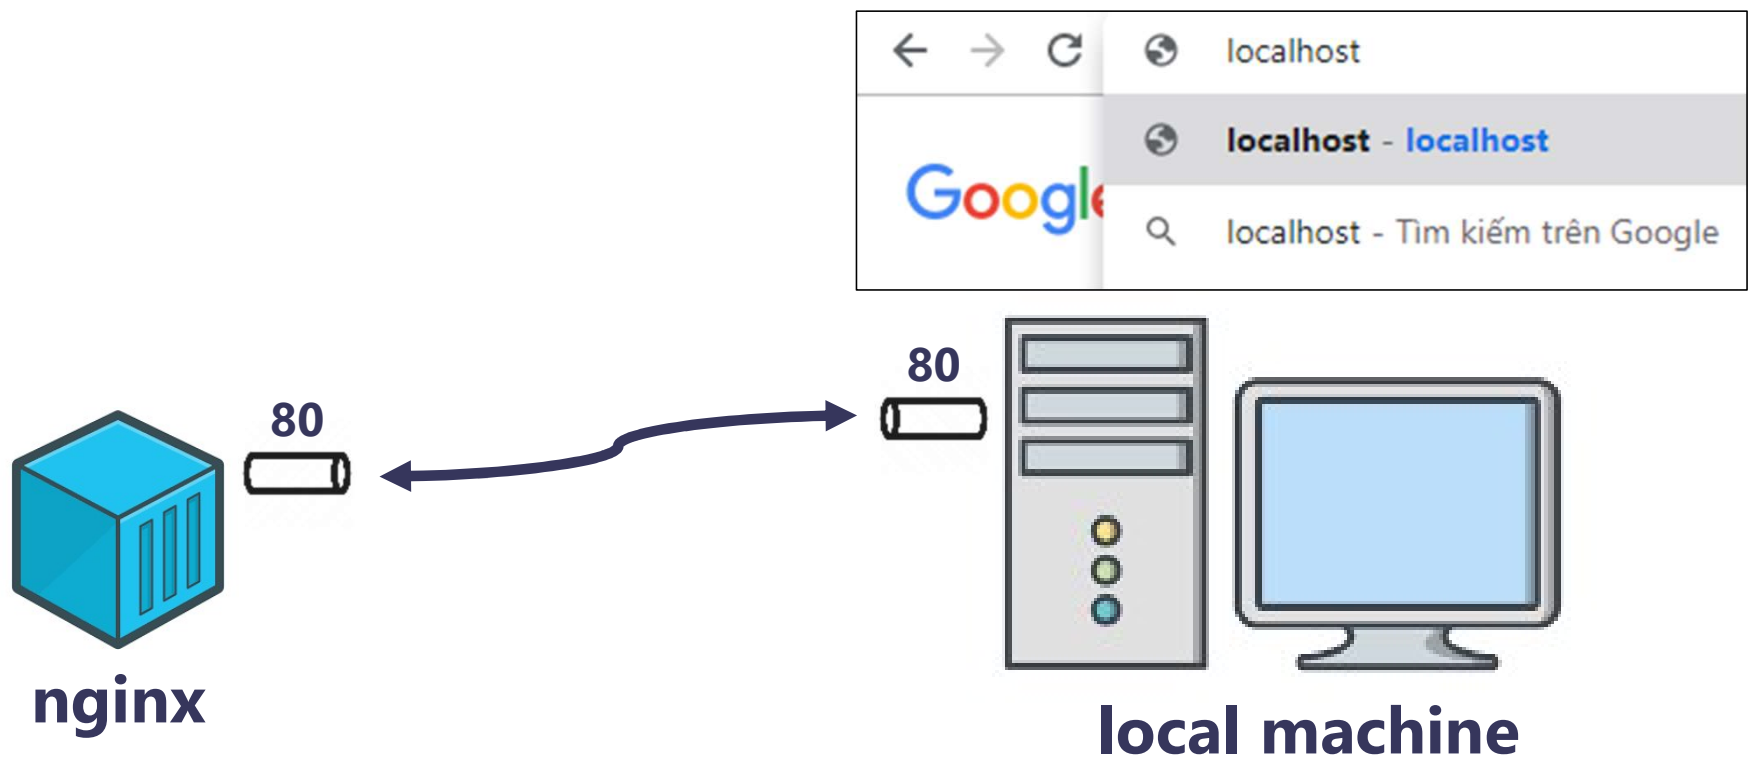


SSH into container:
`docker exec -it ... sh`
`/ # curl localhost`



local machine

3. Port Mapping



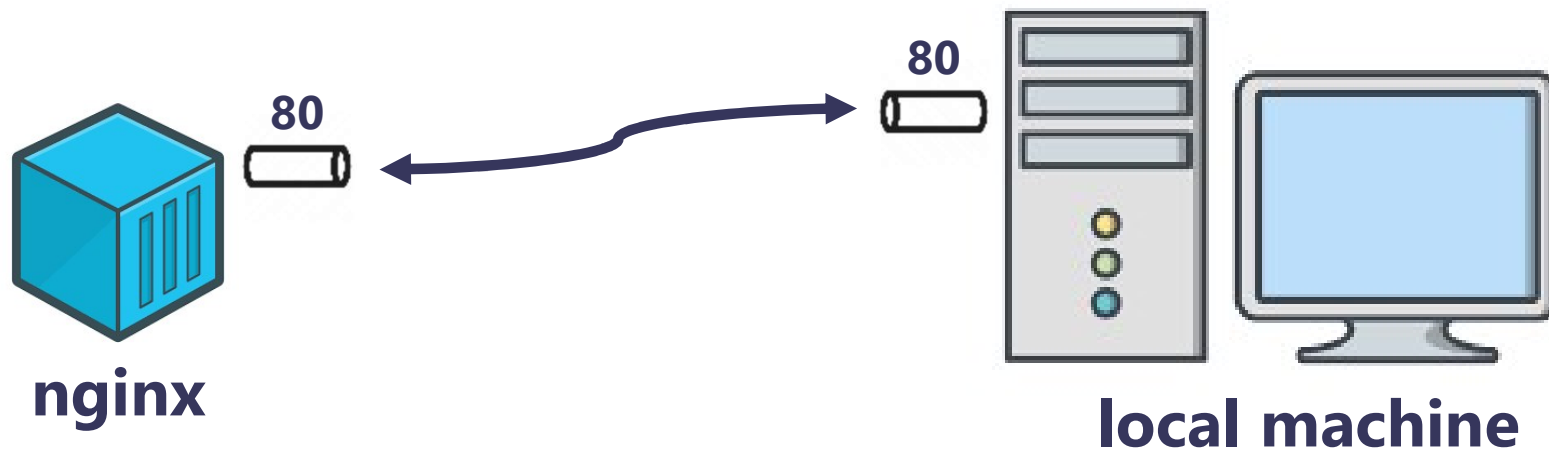
3. Port Mapping

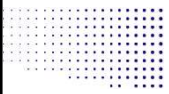
Syntax

- `docker run -p <target_port>:<container_port> ...`

Example

- `docker run -p 80:80 nginx`





3. Logs trace

Syntax

```
docker logs -f <container_id>
```

- -f : keep following the log output

3. Volume – bind mount

Image is immutable. Container is stateless.

```
docker run -p 5432:5432 postgres
```

- How to persist data in postgres? How to make container stateful?
→ Use volume

Volume

- Volume indicates the partition of memory that Docker uses to persist data inside container.



3. Volume – bind mount

Syntax

```
docker volume create [volume_name]
```

```
docker run -v [local_dir/volume]:[container_dir] ...
```

Example

```
docker volume create pgdata
```

```
docker run -v pgdata:/var/lib/postgresql/data -p 5432:5432 postgres
```

```
docker run -v /usr/data:/var/lib/postgresql/data -p 5432:5432 postgres
```



3. Volume – bind mount

Syntax

```
docker volume create [volume_name]
```

```
docker run -v [local_dir/volume]:[container_dir] ...
```

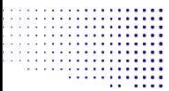
Example

```
docker volume create pgdata
```

```
docker run -v pgdata:/var/lib/postgresql/data -p 5432:5432 postgres
```

```
docker run -v /usr/data:/var/lib/postgresql/data -p 5432:5432 postgres
```

```
docker run -v "C:\users\html":/usr/share/nginx/html -p 80:80 nginx
```



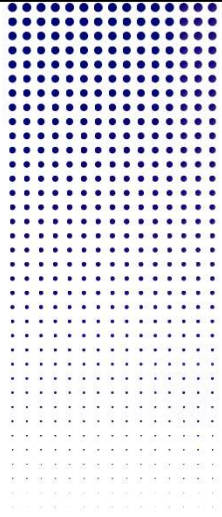
3. Container immutability & statelessness - image

```
docker run -it alpine
```

```
#sh: apk add bash
```

```
docker run -it alpine bash
```

- Error! Modifying container does not affect the original image.
- How to save the desired state of the container?
 - Build our own image



Docker

1

4. How to build my own image?
Dockerfile AZ

4. Dockerfile

- Dockerfile is a template that allows us to instruct Docker to build our own image step by step.
- Back to our example with alpine having bash

Syntax

- `docker build -t <image_name>:<tag> .`

```
FROM alpine:latest
RUN apk add bash
CMD ["bash"]
```

Dockerfile

Build context

Hook command



4. Build context - `.dockerignore`

- “Build context” refers to the folder that contains Dockerfile. All contents inside that folder is called build context.
- Docker client will send build context to Docker Daemon inside Linux machine to build the image.
- Be careful: remember put all necessary resources (files, images...) into the build context, as well as remove all unnecessary things, in order to utilize the size of request payload.

`.dockerignore`

- When sending build context to Docker Daemon, it will ignore files and folders listed in the `.dockerignore` – just the same as `.gitignore`



4. Demo

```
docker build -t my-alpine .
```

```
docker tag <image_name> <new_name>
```

```
FROM alpine:latest
```

```
RUN apk add bash
```

```
CMD ["bash"]
```

Dockerfile

4. Dockerfile keywords

FROM <image>

RUN <command>

WORKDIR <directory>

COPY <src> <dest>

ADD <src/URL> <dest>

EXPOSE <port>

CMD command argument1 argument2...

CMD ["command", "argument1", "argument2", ...]

Shell form

Exec form
(prevent shell injection – recommend)

4. Dockerfile: samples – backend (Java)

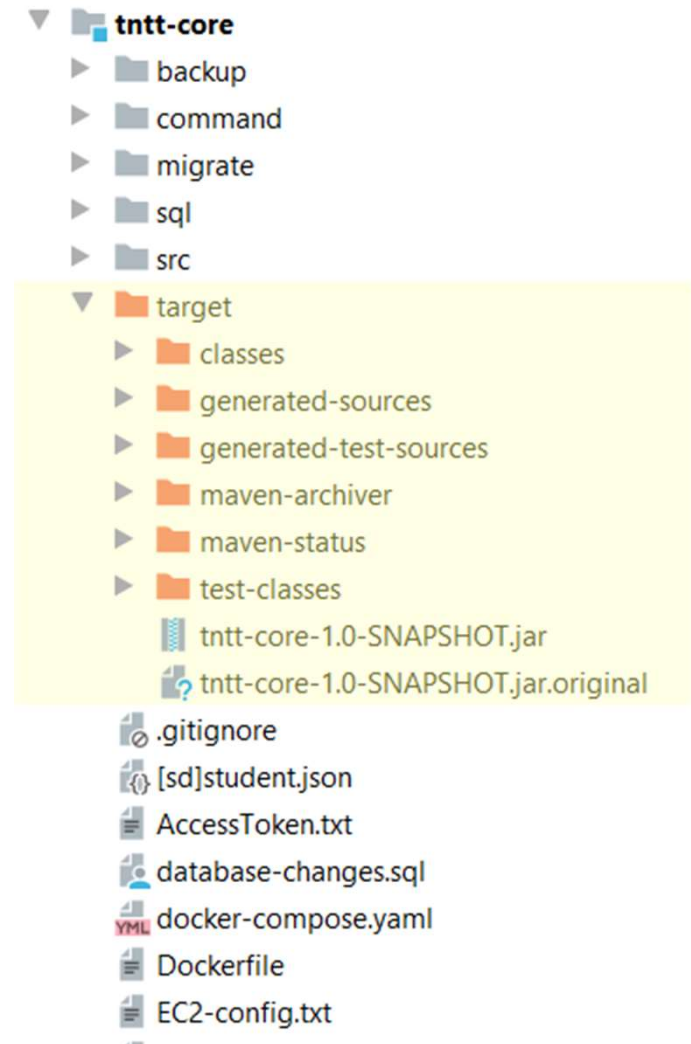
```
FROM openjdk:8-jre
```

```
EXPOSE 8080
```

```
WORKDIR /app
```

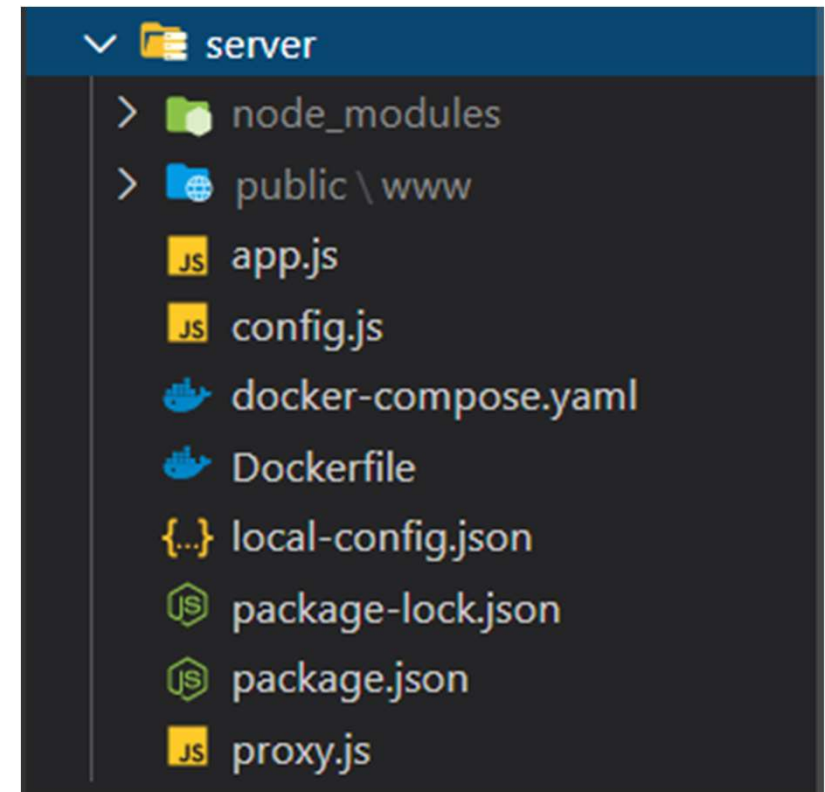
```
COPY ./target/app-1.0-SNAPSHOT.jar .
```

```
CMD ["java", "-jar", "app-1.0-SNAPSHOT.jar"]
```



4. Dockerfile: samples – frontend (ExpressJS)

```
FROM node:alpine
WORKDIR /server
COPY ./*.json ./
COPY ./node_modules ./node_modules
COPY ./*.js ./
COPY ./public ./public
CMD ["node", "app.js"]
```

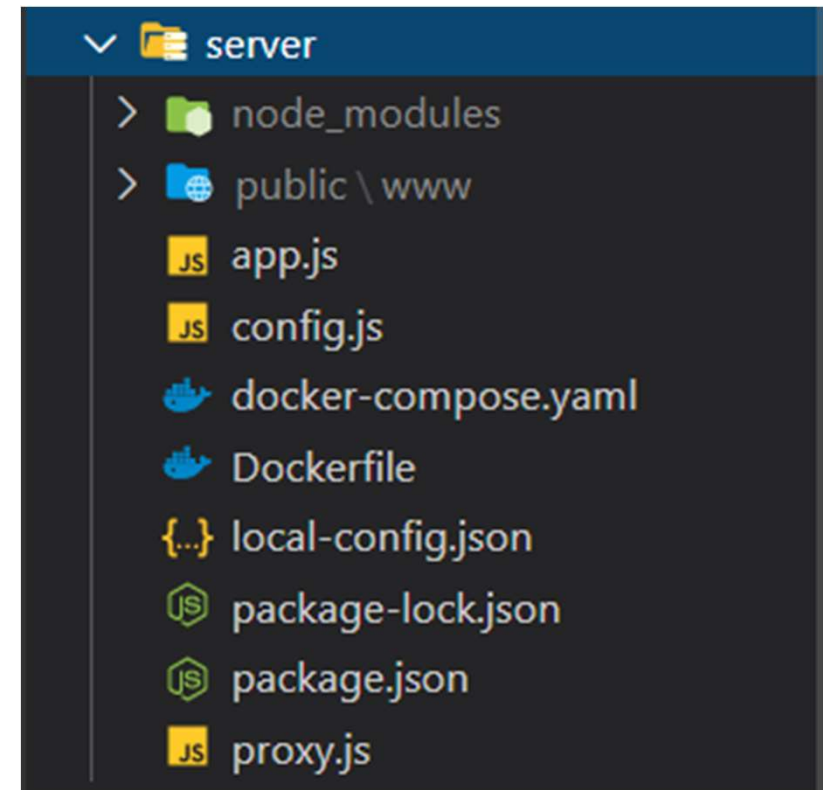


4. Dockerfile: samples – frontend (ExpressJS)

```
FROM node:alpine
WORKDIR /server
COPY ./*.json ./
COPY ./node_modules ./node_modules
COPY ./*.js ./
COPY ./public ./public
CMD ["node", "app.js"]
```

Why do we have to separate multiple commands?

→ Cache upon difference of ***frequency of change of build context***



4. Dockerfile: Utilizing size by layers

Build process

- Image is an ordered set of layers. Each layer is according to a command in Dockerfile.
- With each command, Docker will create a temporary container from previous layer, then execute command inside that container, take a snapshot of it (commit) into a new layer, and finally remove the temporary container which is no longer needed.
- In order to minimize the size of final image, we can combine multiple command using `&&` and clear cached data after those commands.

```
RUN apt-get update \  
    && apt-get install <app> \  
    && rm -rf /var/lib/apt/lists/*
```



4. Dockerfile: Utilizing size by layers

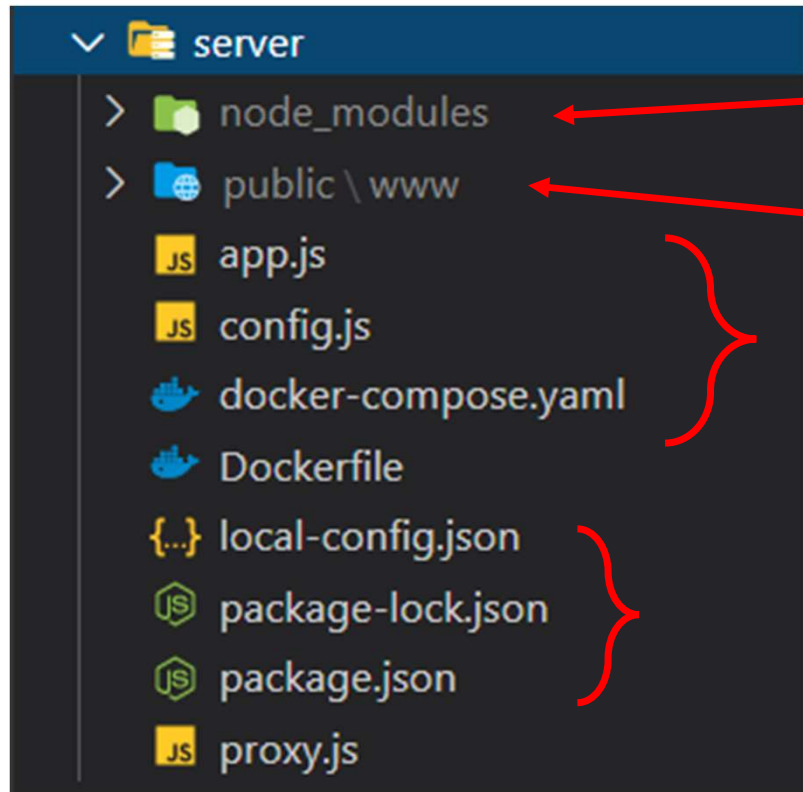
Pushing image

- Login into the registry with **docker login**
- Push image: **docker push <image>:<tag>**

Image cache

- Docker caches layers that remain unchanged (by computing hash value of output of the command). Only changed layers will be pushed.
- To clear cache, use argument **--no-cache** when building image
- Good to remember:
 - It is nice if we could arrange the commands reasonably to minimize the changes made to the image, so that minimize the request payload sent to Docker registry when pushing image.

4. Dockerfile: Utilizing size by layers

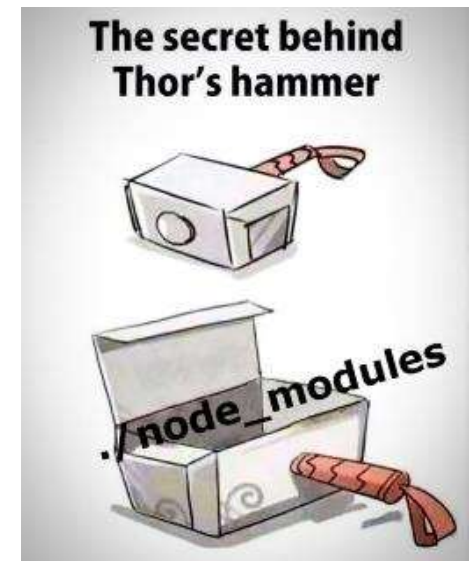


Libraries of ExpressJS

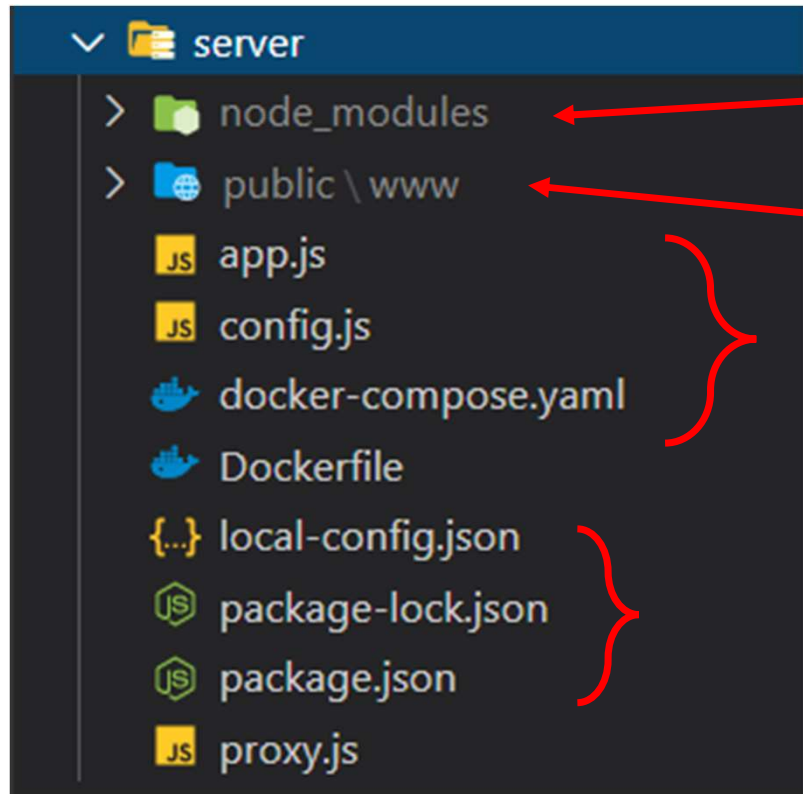
Static HTML files

ExpressJS source code

Manifest files



4. Dockerfile: Utilizing size by layers



1. Libraries of ExpressJS

3. Static HTML files

2. ExpressJS source code

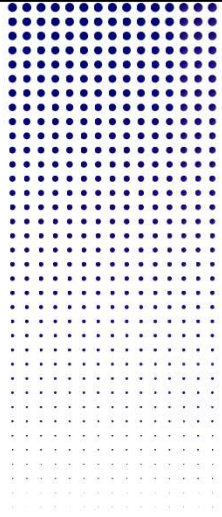
1. Manifest files

4. Dockerfile: Utilizing size by layers

```
FROM node:alpine
WORKDIR /server
COPY ./*.json ./
COPY ./node_modules ./node_modules
COPY ./*.js ./
COPY ./public ./public
CMD ["node", "app.js"]
```

```
FROM node:alpine
WORKDIR /server
COPY ./public ./public
COPY ./*.json ./
COPY ./node_modules ./node_modules
COPY ./*.js ./
CMD ["node", "app.js"]
```

Which one is better?



2

Docker compose



Agenda

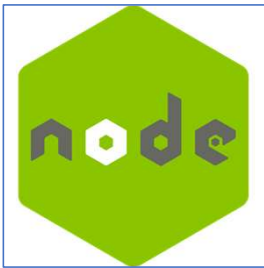
1. Docker

2. **Docker compose**

1. *Multiple containers: configurations and communication*
2. *docker-compose.yml*
3. *Build image easily with compose*
4. *Run container easily with compose*

3. Docker Swarm

1. Multiple containers



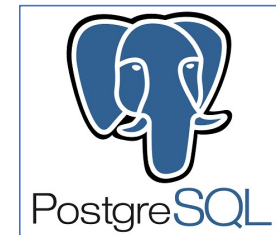
Frontend

port: 80
envs: [...]



Backend

port: 8080
envs: [...]



Database

port: 5432
volume: /usr/data
envs: [...]

1. docker-compose

docker-compose.yml

Frontend

port: 80
envs: [...]

Backend

port: 8080
envs: [...]

Database

port: 5432
volume: /usr/data
envs: [...]

comment # This is comment

object: yaml

```
title: Vừa Nhắm Mắt Vừa Mở Cửa Sổ  
author:  
  name: Nguyễn Ngọc Thuần  
  birthYear: 1972
```

object: json

```
{  
  title: "Vừa Nhắm Mắt Vừa Mở Cửa Sổ",  
  author: {  
    name: "Nguyễn Ngọc Thuần"  
    birthYear: 1972  
  }  
}
```

array: yaml

```
publishers:  
  - name: NXB Trẻ  
    year: 2004  
  - name: NXB Văn học  
    year: 2010
```

array: json

```
{  
  publishers: [  
    { name: "NXB Trẻ", year: 2004 },  
    { name: "NXB Văn học", year: 2010 }  
  ]  
}
```



1. docker-compose.yml

```
version: '3'
services:
  pg:
    image: postgres:9.6-alpine
    ports:
      - 5432:5432
    volumes:
      - pgdata:/var/lib/postgresql/data
    environment:
      POSTGRES_DB: postgres
      POSTGRES_USER: postgres
      POSTGRES_PASSWORD: postgres

  frontend:
    image: nambach/frontend:latest
    build:
      context: .

volumes:
  pgdata:
```

2. Build image easily with compose

Place docker-compose.yaml in the same folder of Dockerfile and build.

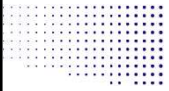
Syntax

```
docker-compose build <service_name>
```

Example

```
docker-compose build frontend
```

```
version: '3'
services:
  frontend:
    image: nambach/frontend:latest
    build:
      context: .
```



3. Run container easily with compose

Syntax

```
docker-compose up
```

```
docker-compose up <service_name>
```

```
docker-compose up -d <service_name>
```

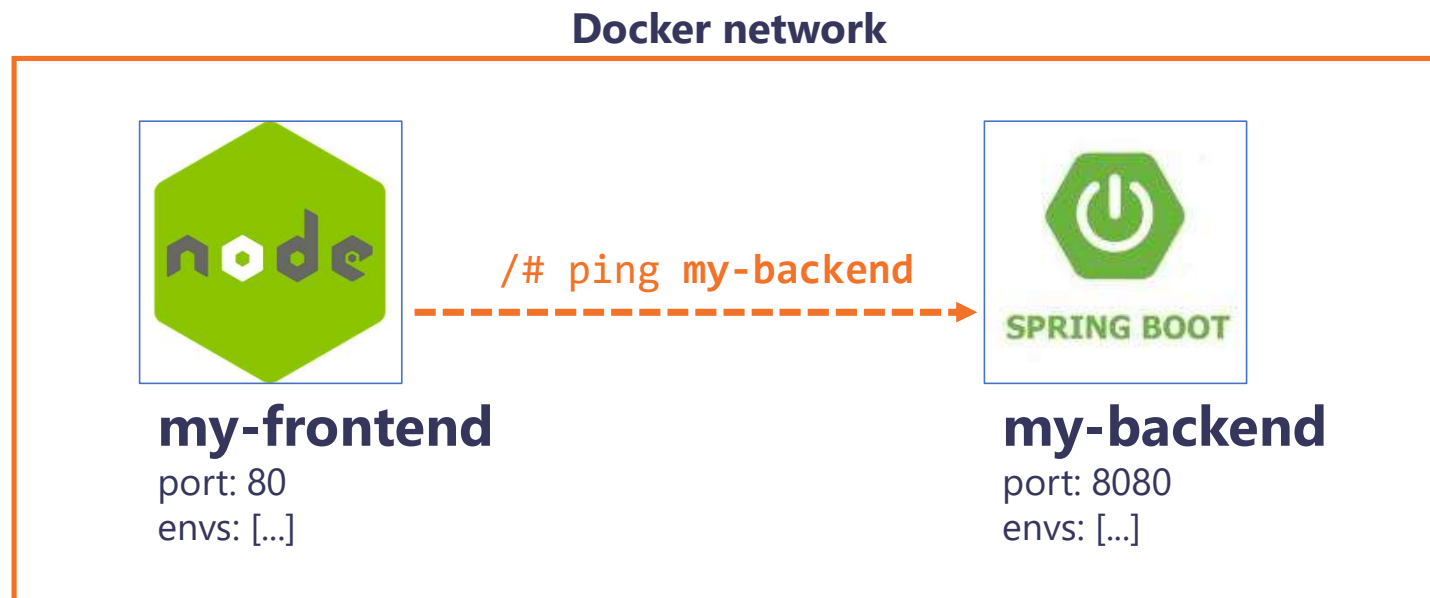
```
docker-compose logs -f <service_name>
```

```
docker-compose stop <service_name>
```

```
docker-compose down
```

4. Network

When attaching containers into same network, thanks to the built-in network resolution, Docker can resolve service's name into its actual IP. So that containers can now communicate with each other through service names instead of IPs.





Demo

Syntax

```
docker-compose up
```

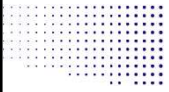
```
docker-compose up <service_name>
```

```
docker-compose up -d <service_name>
```

```
docker-compose logs -f <service_name>
```

```
docker-compose stop <service_name>
```

```
docker-compose down
```



Agenda

1. Docker
2. **Docker compose**
3. Docker Swarm



THANK YOU